

History-Independent Load Balancing



Michael A. Bender
Stony Brook University



Bill Kuszmaul
CMU



Elaine Shi
CMU



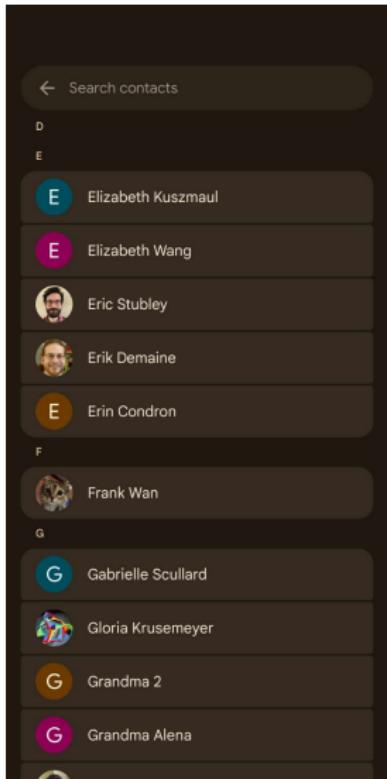
Rose Silver
CMU

HISTORY INDEPENDENT DATA STRUCTURES

History Independence: “If an adversary were to see the state of the data structure, they would learn only the current set of elements, and nothing else about the history of past operations.”

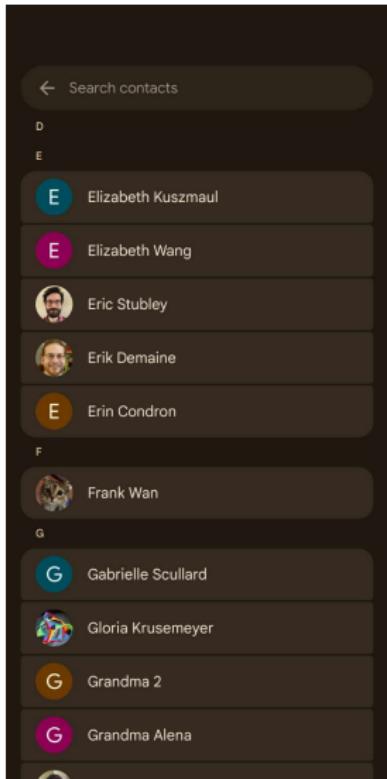
[Micciancio '97], [Naor, Teague '01]

HISTORY VS CONTENT



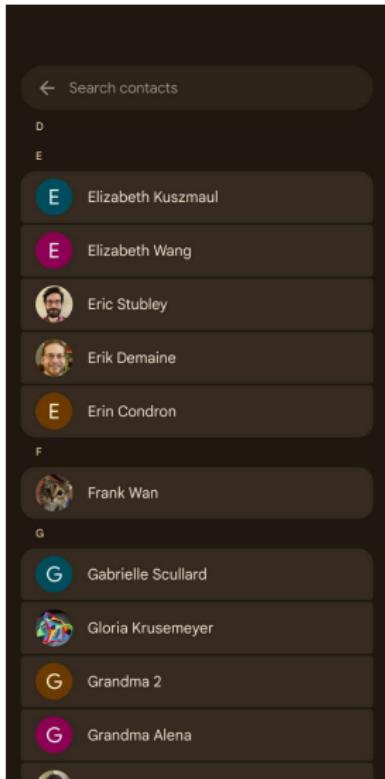
- ▶ If someone hacks my phone, they can learn my contacts list.

HISTORY VS CONTENT



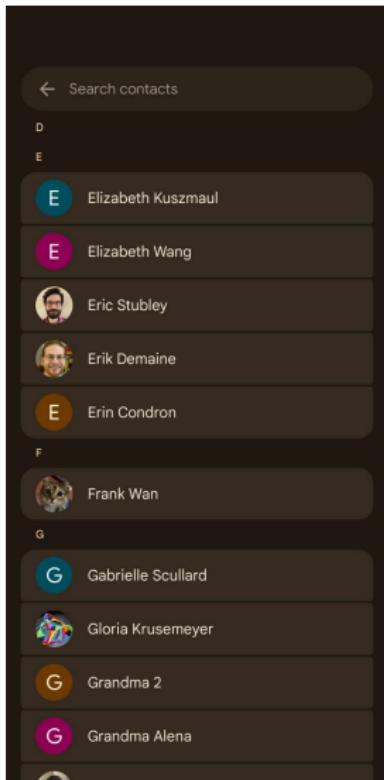
- ▶ If someone hacks my phone, they can learn my contacts list.
- ▶ But can they learn who my contacts were in the past?

HISTORY VS CONTENT



- ▶ If someone hacks my phone, they can learn my contacts list.
- ▶ But can they learn who my contacts were in the past?
- ▶ What about the order in which contacts were added?

HISTORY VS CONTENT



- ▶ If someone hacks my phone, they can learn my contacts list.
- ▶ But can they learn who my contacts were in the past?
- ▶ What about the order in which contacts were added?
- ▶ A history independent data structure protects this kind of information.

HISTORY INDEPENDENT IS A SECURITY GUARANTEE

History Independence: “If an adversary were to see the state of the data structure, they would learn only the current set of elements, and nothing else about the history of past operations.”

[Micciancio '97], [Naor, Teague '01]

HISTORY INDEPENDENT IS A SECURITY GUARANTEE

History Independence: “If an adversary were to see the state of the data structure, they would learn only the current set of elements, and nothing else about the history of past operations.”

[Micciancio '97], [Naor, Teague '01]

Lost of successes: Hash tables, trees, memory allocation, PMAs, graph algorithms, B-trees, cache-oblivious data structures...

[Micciancio '97], [Naor, Teague '01], [Buchbinder, Petrank '03], [Molnar, Kohno, Sastry, Wagner '06], [Blelloch, Golovin '07], [Moran, Naor, Segev '07] [Naor, Segev, Wieder '08], [Golovin '08 '09 '10], [Tzouramanis '12], [Bajaj, Sion '13] [Bajaj, Chakrabati, Sion '15], [Roche, Aviv, Choi '15], [Bender, Berry, Johnson, Kroeger, McCauley, Phillips, Simon, Singh, Zage '16], ...

HISTORY INDEPENDENT IS A SECURITY GUARANTEE

History Independence: “If an adversary were to see the state of the data structure, they would learn only the current set of elements, and nothing else about the history of past operations.”

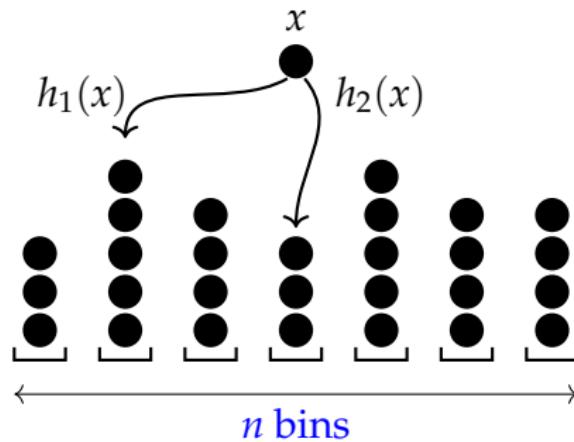
[Micciancio '97], [Naor, Teague '01]

Lost of successes: Hash tables, trees, memory allocation, PMAs, graph algorithms, B-trees, cache-oblivious data structures...

[Micciancio '97], [Naor, Teague '01], [Buchbinder, Petrank '03], [Molnar, Kohno, Sastry, Wagner '06], [Blelloch, Golovin '07], [Moran, Naor, Segev '07] [Naor, Segev, Wieder '08], [Golovin '08 '09 '10], [Tzouramanis '12], [Bajaj, Sion '13] [Bajaj, Chakrabati, Sion '15], [Roche, Aviv, Choi '15], [Bender, Berry, Johnson, Kroeger, McCauley, Phillips, Simon, Singh, Zage '16], ...

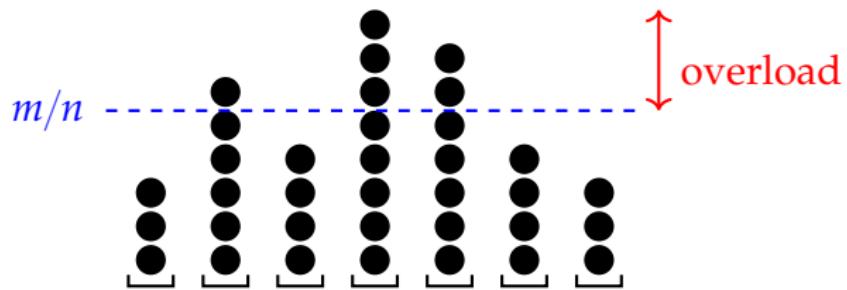
But... some very basic questions also remain open.

TWO-CHOICE LOAD BALANCING



- ▶ Balls are **inserted/deleted**, with up to m present at a time.
- ▶ Each ball has two random bins where it can go.
- ▶ We must maintain a valid assignment of balls to bins.

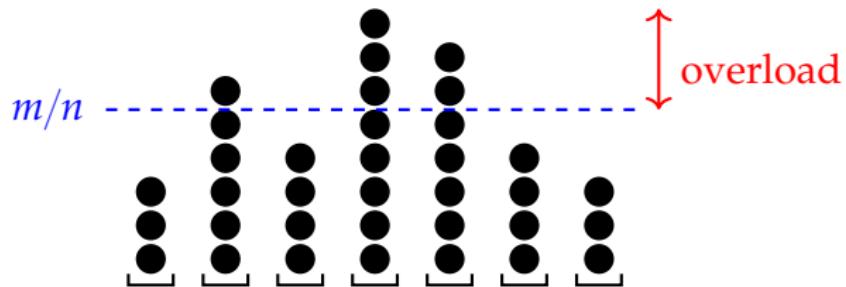
TWO GOALS



Minimize Overload:

The amount by which the fullest bin exceeds m/n is small.

TWO GOALS



Minimize Overload:

The amount by which the fullest bin exceeds m/n is small.

Minimize Recourse:

On any given insertion/deletion, the number of balls moved around is small.

THIS PAPER

Question: Does there exist a history-independent solution with small recourse and overload?

THIS PAPER

Question: Does there exist a history-independent solution with small recourse and overload?

Theorem: There exists a history-independent solution with:

- ▶ Overload $O(1)$, with high probability.
- ▶ Expected recourse $O(\log \log(m/n))$.

WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]
[Frieze and Petti '18] ...

WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]
[Frieze and Petti '18] ...

But the fully dynamic case has remained largely open.

[Vöcking '99] [Dietzfelbinger and Weidling '07] [Bender, Conway, Farach-Colton, Kuszmaul, Tagliavini '21] [Bansal, Kuszmaul '22] ...

WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]
[Frieze and Petti '18] ...

But the fully dynamic case has remained largely open.

[Vöcking '99] [Dietzfelbinger and Weidling '07] [Bender, Conway, Farach-Colton, Kuszmaul, Tagliavini '21] [Bansal, Kuszmaul '22] ...

Open Question:

Is there a **fully dynamic** solution with **recourse $o(m/n)$** and **overload $O(\log \log n)$** ?

WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]
[Frieze and Petti '18] ...

But the fully dynamic case has remained largely open.

[Vöcking '99] [Dietzfelbinger and Weidling '07] [Bender, Conway, Farach-Colton, Kuszmaul, Tagliavini '21] [Bansal, Kuszmaul '22] ...

Open Question:

Is there a **fully dynamic** solution with **recourse $o(m/n)$** and **overload $O(\log \log n)$** ?

Answer:

Yes! We get **recourse $O(\log \log(m/n))$** and **overload $O(1)$** .

THIS PAPER

Question: Does there exist a history-independent solution with small recourse and overload?

Theorem: There exists a history-independent solution with:

- ▶ Overload $O(1)$, with high probability.
- ▶ Expected recourse $O(\log \log(m/n))$.

THIS PAPER

Question: Does there exist a history-independent solution with small recourse and overload?

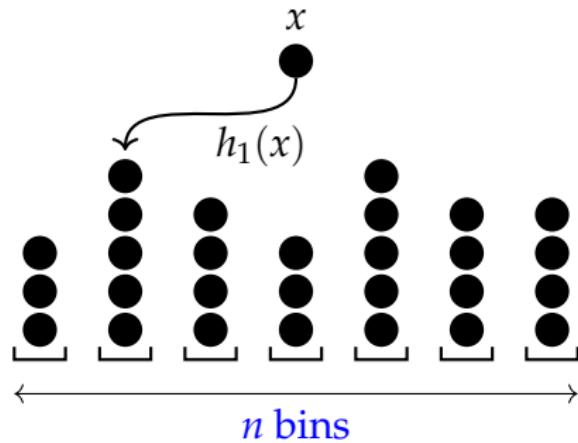
Theorem: There exists a history-independent solution with:

- ▶ Overload $O(1)$, with high probability.
- ▶ Expected recourse $O(\log \log(m/n))$.

Rest of Talk: A simple history-independent algorithm with overload $O(\log \log n)$ and expected recourse $O(m/n)$.

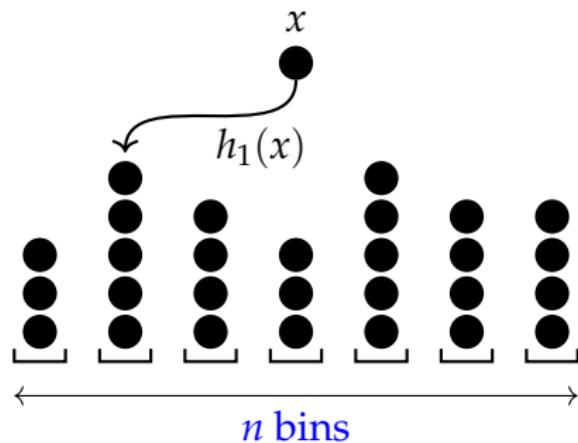
WARMUP 1: THE SINGLE-CHOICE STRATEGY

To insert a ball x , just put it in bin $h_1(x)$:



WARMUP 1: THE SINGLE-CHOICE STRATEGY

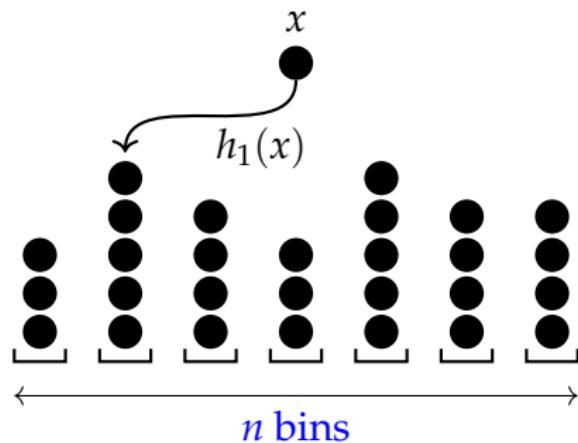
To insert a ball x , just put it in bin $h_1(x)$:



- ▶ This is history-independent ✓

WARMUP 1: THE SINGLE-CHOICE STRATEGY

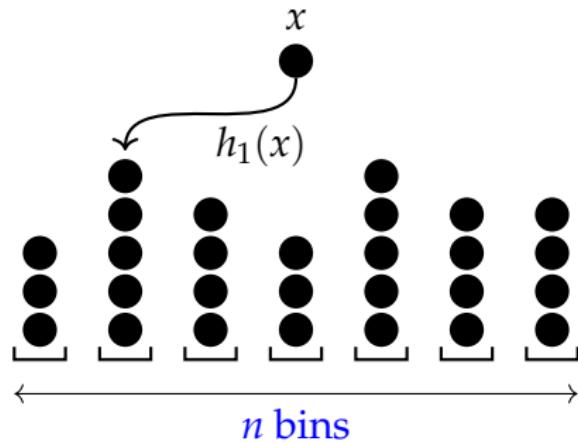
To insert a ball x , just put it in bin $h_1(x)$:



- ▶ This is history-independent ✓
- ▶ The recourse is 0 ✓

WARMUP 1: THE SINGLE-CHOICE STRATEGY

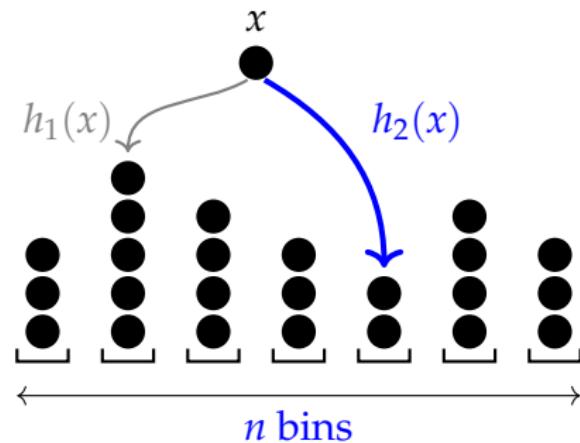
To insert a ball x , just put it in bin $h_1(x)$:



- ▶ This is history-independent ✓
- ▶ The recourse is 0 ✓
- ▶ But... the overload is huge, roughly $\sqrt{m/n}$ ✗

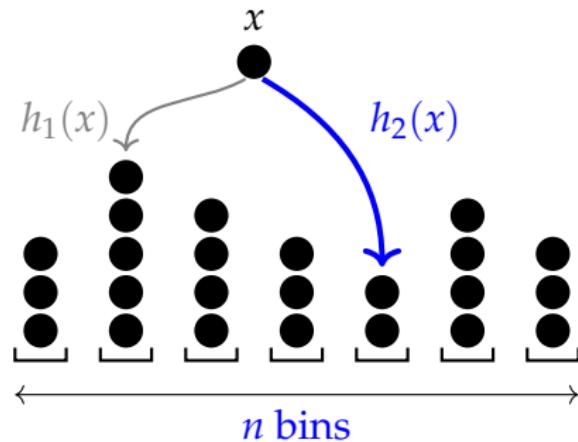
WARMUP 2: GREEDY INSERTIONS

To insert a ball x , put it in the **emptier** of its choices:



WARMUP 2: GREEDY INSERTIONS

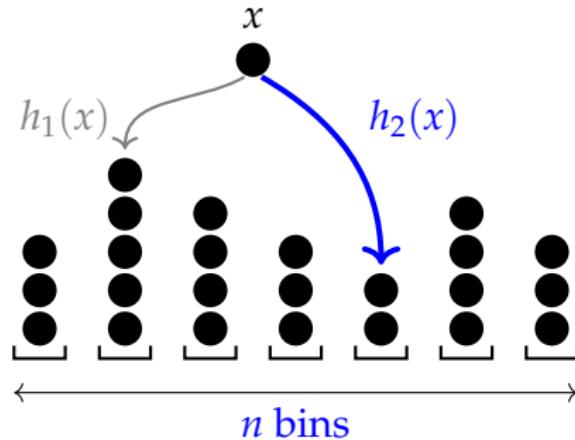
To insert a ball x , put it in the **emptier** of its choices:



- ▶ This is **not** history-independent \times

WARMUP 2: GREEDY INSERTIONS

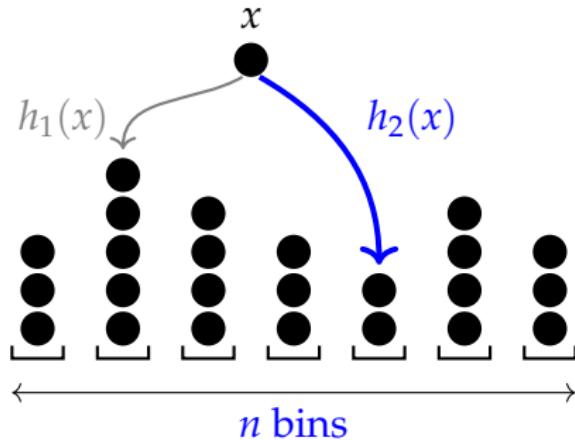
To insert a ball x , put it in the **emptier** of its choices:



- ▶ This is **not** history-independent ✗
- ▶ The recourse is 0 ✓

WARMUP 2: GREEDY INSERTIONS

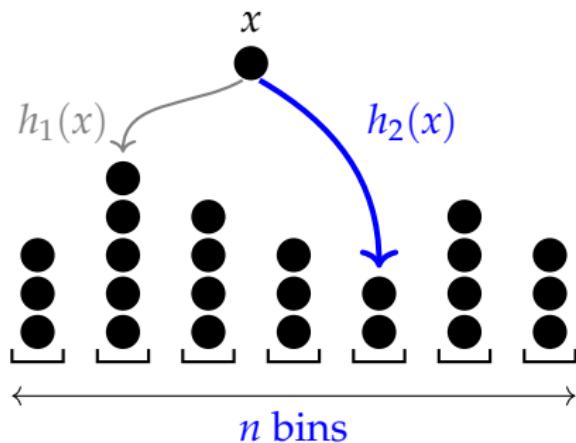
To insert a ball x , put it in the **emptier** of its choices:



- ▶ This is **not** history-independent ✗
- ▶ The recourse is 0 ✓
- ▶ In the insertion-only case, the overload is $O(\log \log n)$ ✓

[Azar, Broder, Karlin and Upfal '94]

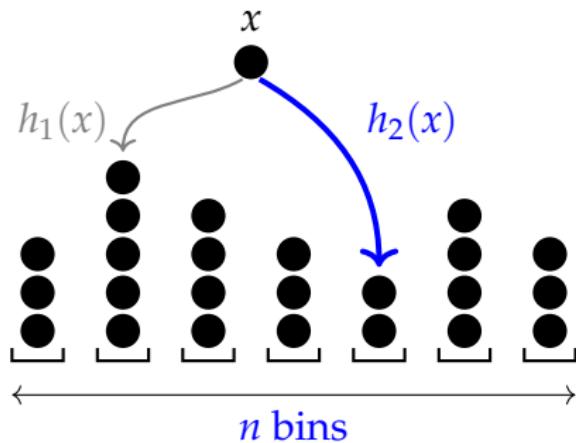
A SIMPLE HISTORY-INDEPENDENT ALGORITHM



Given a set S of balls, define $\text{Greedy}(S)$ as:

- ▶ Start with empty bins.
- ▶ Sort the balls in S to get a sequence x_1, x_2, \dots
- ▶ Insert x_1, x_2, \dots using the greedy algorithm.

A SIMPLE HISTORY-INDEPENDENT ALGORITHM

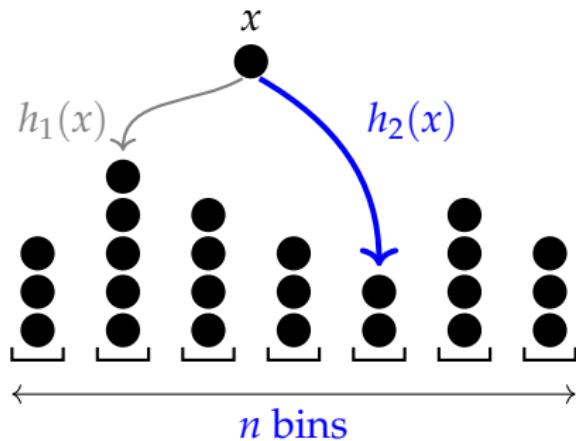


Given a set S of balls, define $\text{Greedy}(S)$ as:

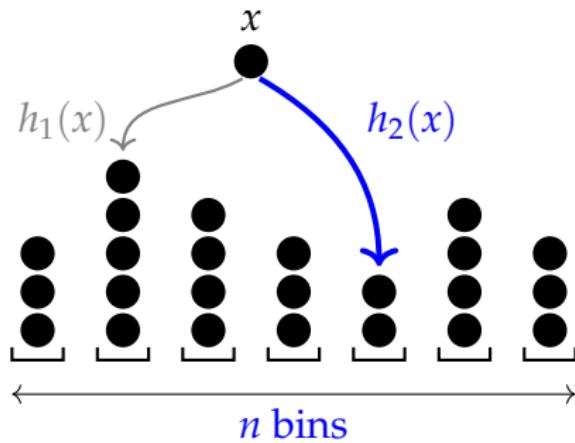
- ▶ Start with empty bins.
- ▶ Sort the balls in S to get a sequence x_1, x_2, \dots
- ▶ Insert x_1, x_2, \dots using the greedy algorithm.

A History-Independent Algorithm: At any given moment, if S is the current set of balls, use allocation $\text{Greedy}(S)$.

ANALYZING HISTORY-INDEPENDENT GREEDY

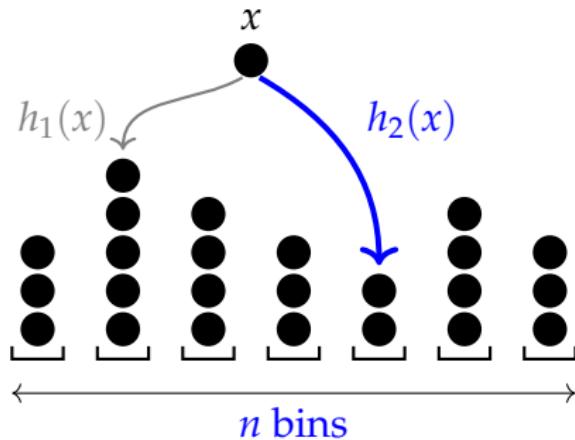


ANALYZING HISTORY-INDEPENDENT GREEDY



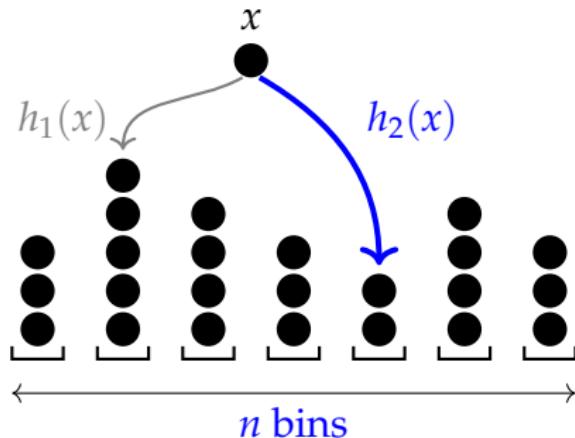
- ▶ The algorithm is history independent ✓

ANALYZING HISTORY-INDEPENDENT GREEDY



- ▶ The algorithm is history independent ✓
- ▶ The overload is $O(\log \log n)$ ✓

ANALYZING HISTORY-INDEPENDENT GREEDY



- ▶ The algorithm is history independent ✓
- ▶ The overload is $O(\log \log n)$ ✓
- ▶ But what about the recourse?

ANALYZING THE RE COURSE



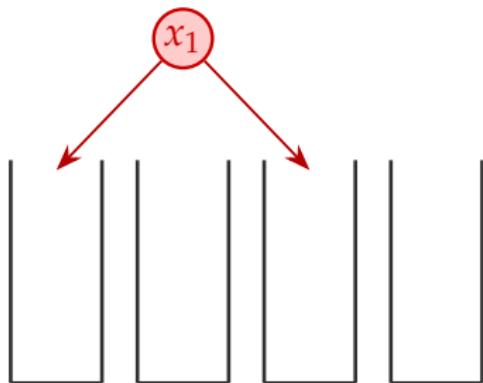
Computing **Greedy**(S)



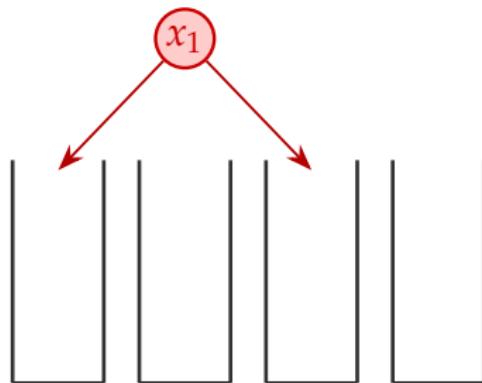
Computing **Greedy**($S \cup \{x^*\}$)

How does $\text{Greedy}(S)$ change if we add a ball x^* ?

ANALYZING THE RE COURSE



Computing **Greedy**(S)



Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

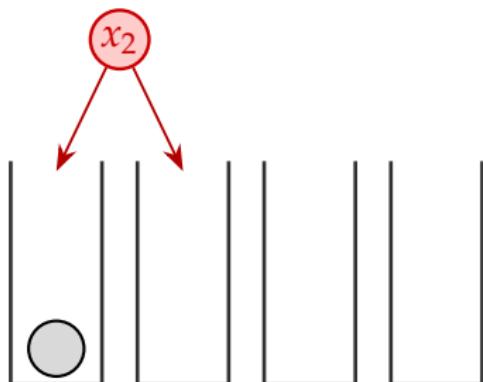


Computing **Greedy**(S)

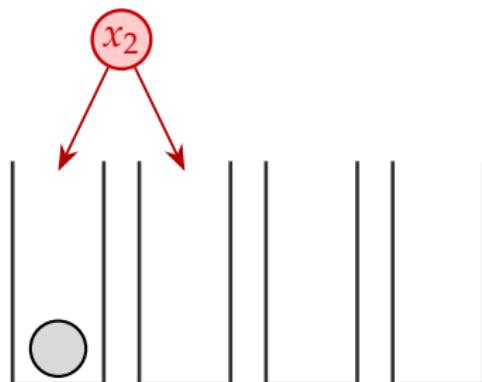


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

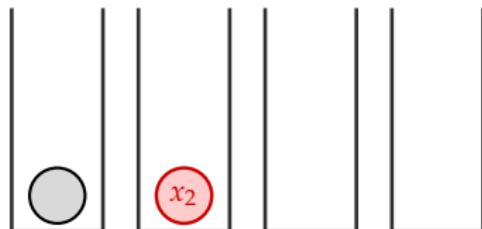


Computing **Greedy**(S)

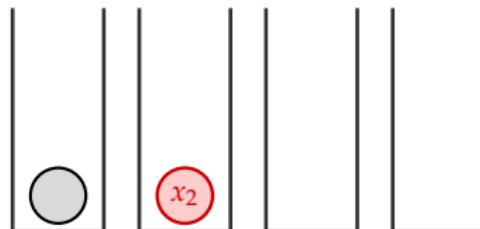


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

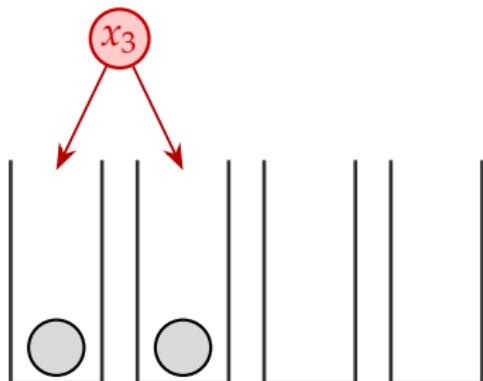


Computing **Greedy**(S)

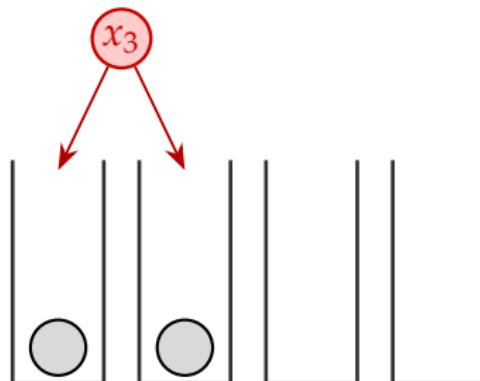


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

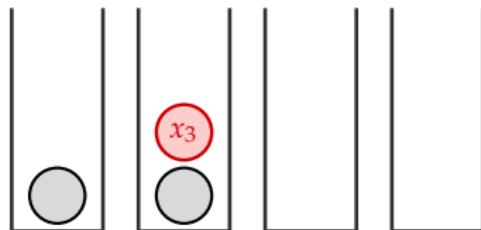


Computing **Greedy**(S)

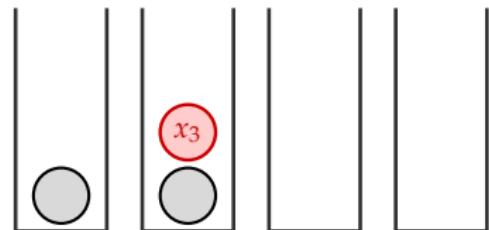


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

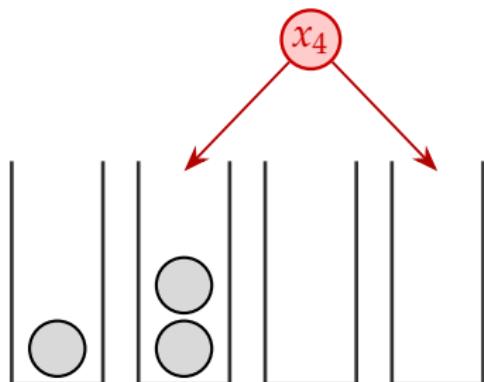


Computing **Greedy**(S)

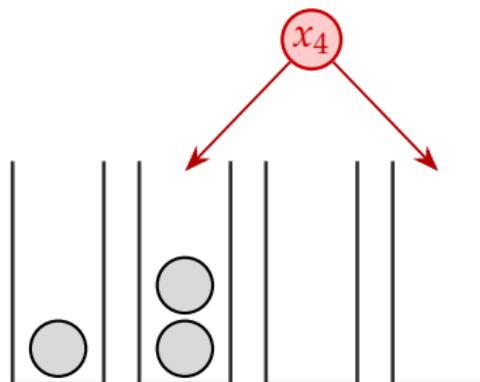


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

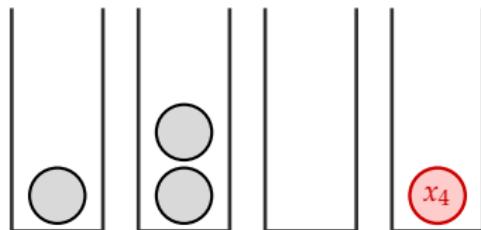


Computing **Greedy**(S)

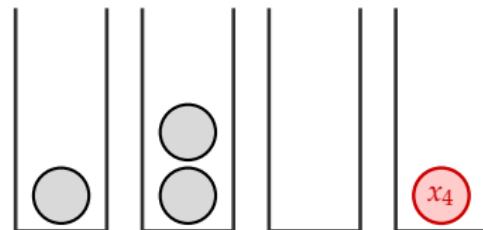


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

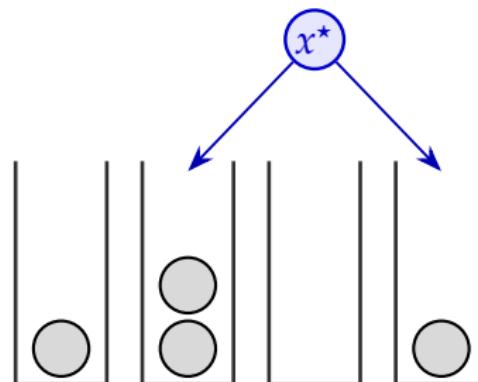
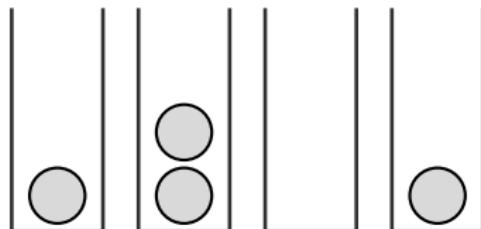


Computing **Greedy**(S)

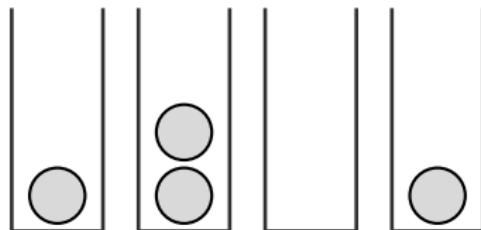


Computing **Greedy**($S \cup \{x^*\}$)

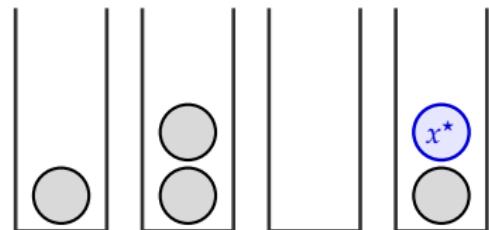
ANALYZING THE RE COURSE



ANALYZING THE RE COURSE

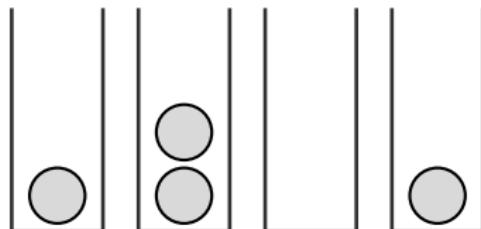


Computing **Greedy**(S)

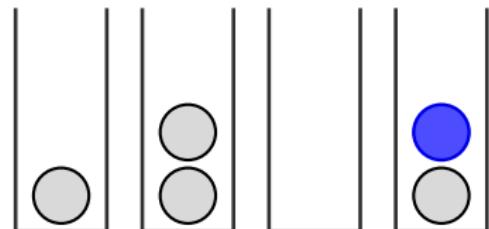


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE

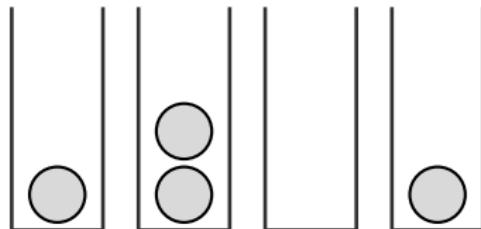


Computing **Greedy**(S)

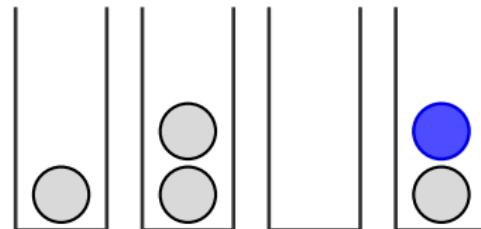


Computing **Greedy**($S \cup \{x^*\}$)

ANALYZING THE RE COURSE



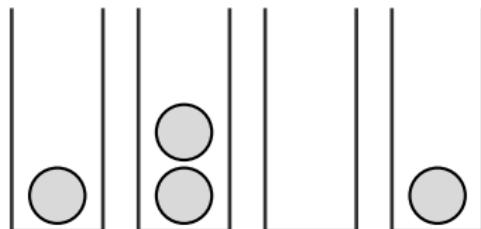
Computing **Greedy**(S)



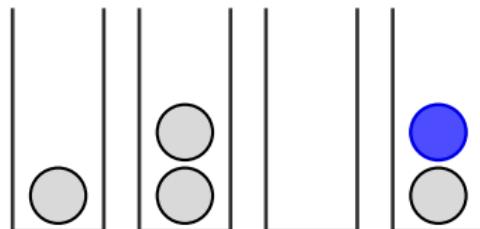
Computing **Greedy**($S \cup \{x^*\}$)

Subsequent balls will experience either:

ANALYZING THE RE COURSE



Computing Greedy(S)

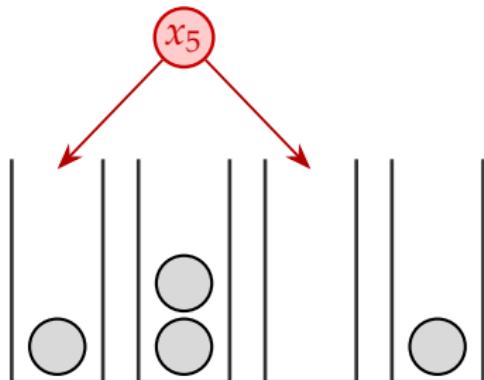


Computing Greedy($S \cup \{x^*\}$)

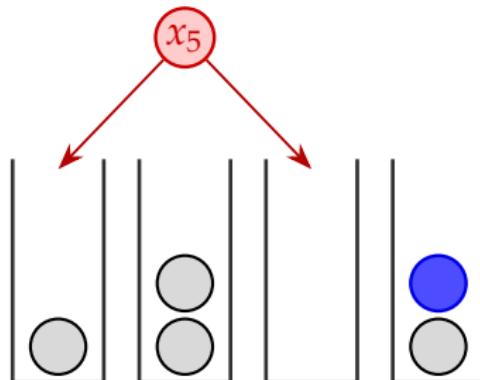
Subsequent balls will experience either:

1. No recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

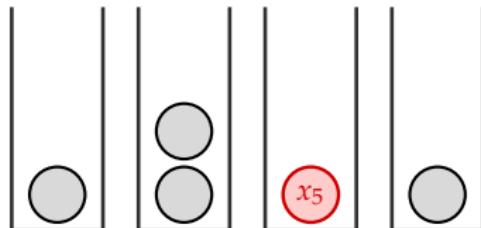


Computing Greedy($S \cup \{x^*\}$)

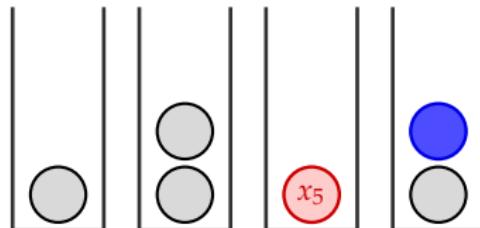
Future insertions will experience either:

1. No recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

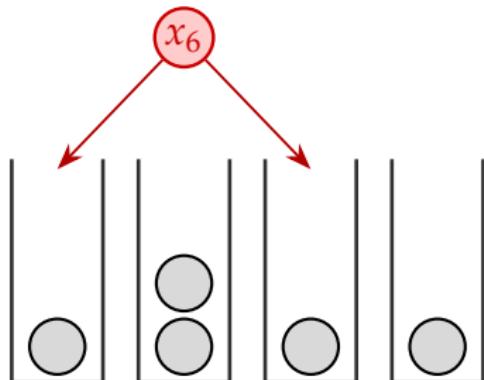


Computing Greedy($S \cup \{x^*\}$)

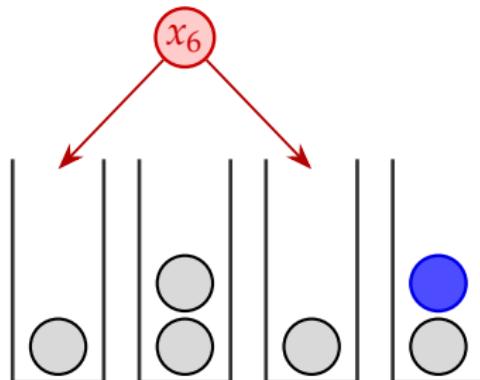
Subsequent balls will experience either:

1. No recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

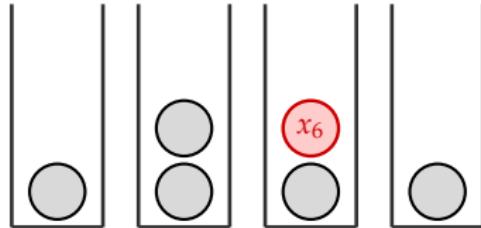


Computing Greedy($S \cup \{x^*\}$)

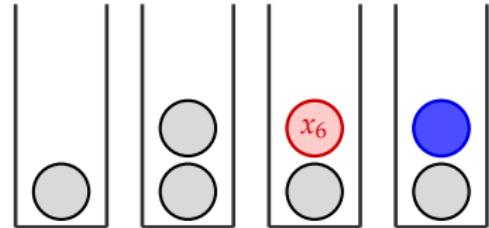
Subsequent balls will experience either:

1. No recourse

ANALYZING THE RE COURSE



Computing **Greedy**(S)

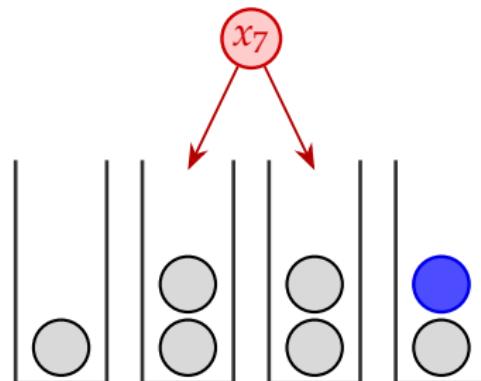
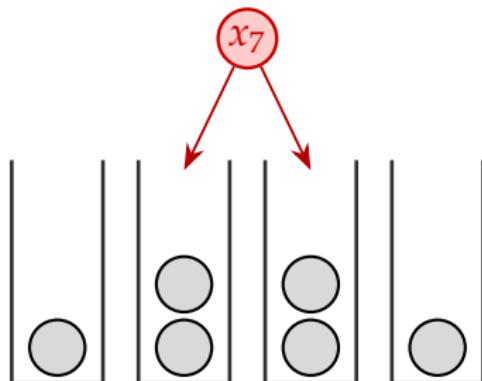


Computing **Greedy**($S \cup \{x^*\}$)

Subsequent balls will experience either:

1. No recourse

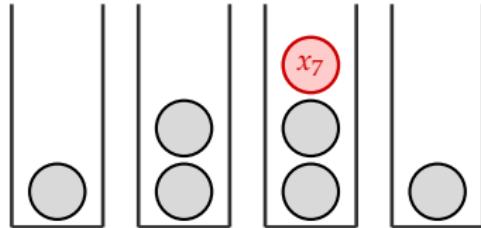
ANALYZING THE RE COURSE



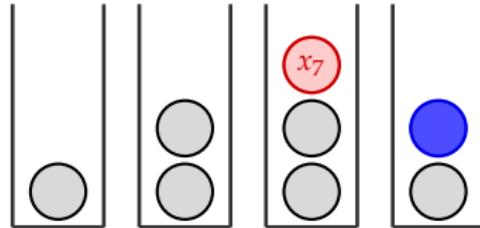
Subsequent balls will experience either:

1. No recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

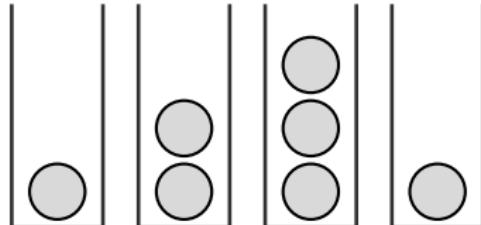


Computing Greedy($S \cup \{x^*\}$)

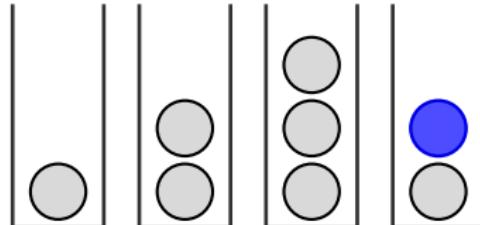
Subsequent balls will experience either:

1. No recourse

ANALYZING THE RE COURSE



Computing **Greedy**(S)

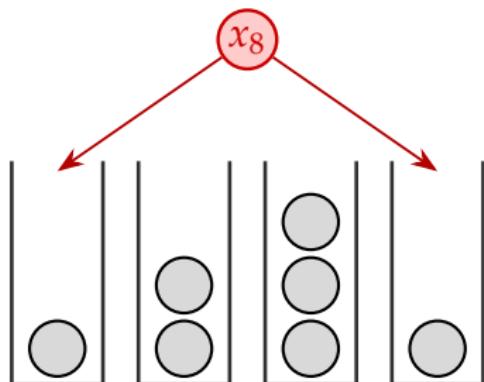


Computing **Greedy**($S \cup \{x^*\}$)

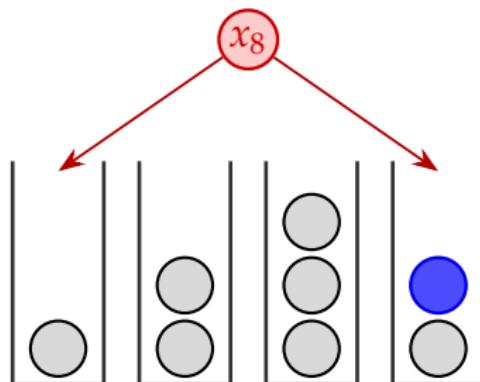
Subsequent balls will experience either:

1. No recourse
2. Recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

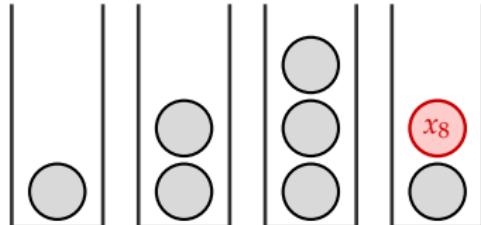


Computing Greedy($S \cup \{x^*\}$)

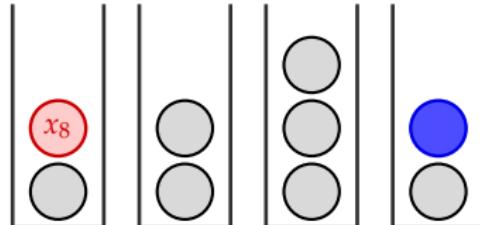
Subsequent balls will experience either:

1. No recourse
2. Recourse

ANALYZING THE RE COURSE



Computing **Greedy**(S)

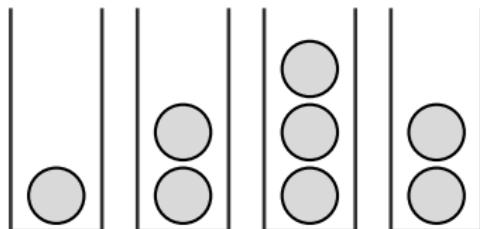


Computing **Greedy**($S \cup \{x^*\}$)

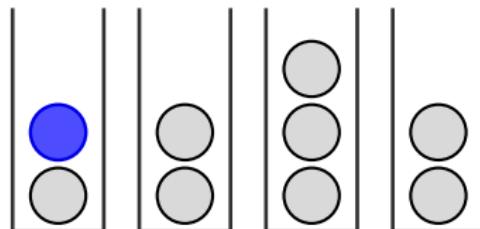
Subsequent balls will experience either:

1. No recourse
2. Recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

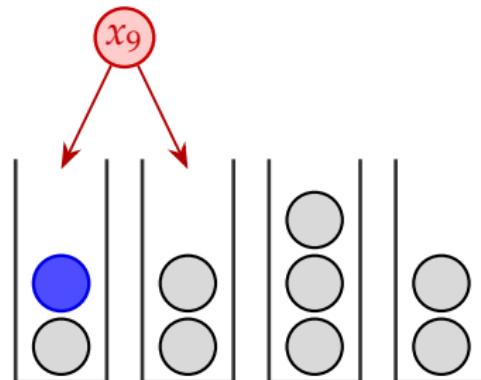
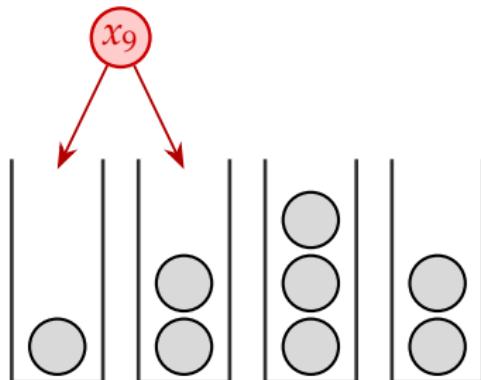


Computing Greedy($S \cup \{x^*\}$)

Subsequent balls will experience either:

1. No recourse
2. Recourse

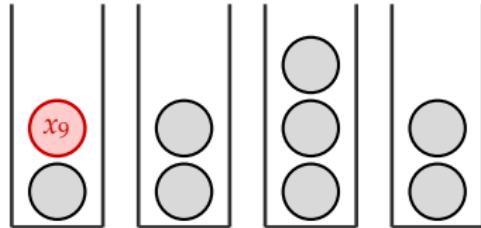
ANALYZING THE RE COURSE



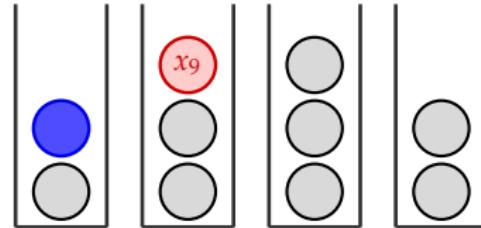
Subsequent balls will experience either:

1. No recourse
2. Recourse

ANALYZING THE RE COURSE



Computing **Greedy**(S)

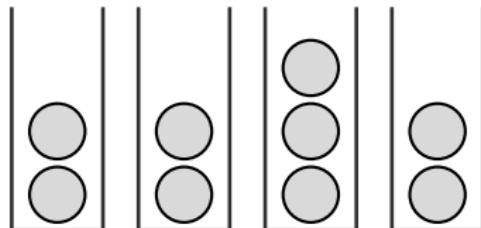


Computing **Greedy**($S \cup \{x^*\}$)

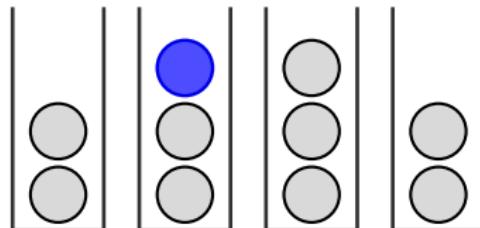
Subsequent balls will experience either:

1. No recourse
2. Recourse

ANALYZING THE RE COURSE



Computing Greedy(S)

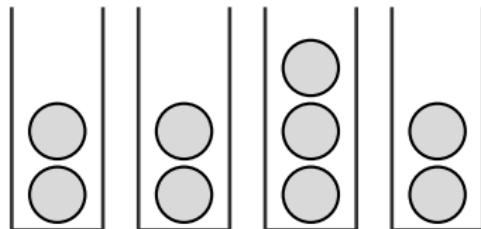


Computing Greedy($S \cup \{x^*\}$)

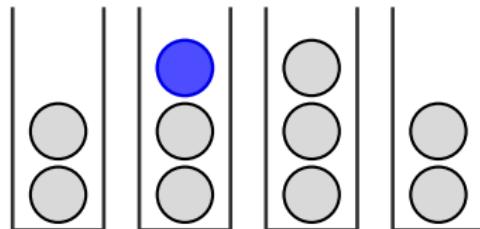
Subsequent balls will experience either:

1. No recourse
2. Recourse

ANALYZING THE RE COURSE



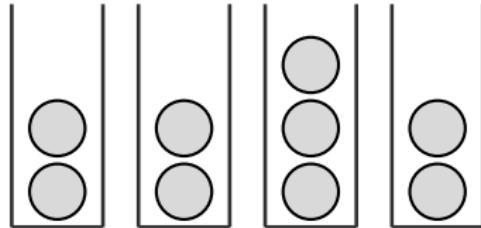
Computing Greedy(S)



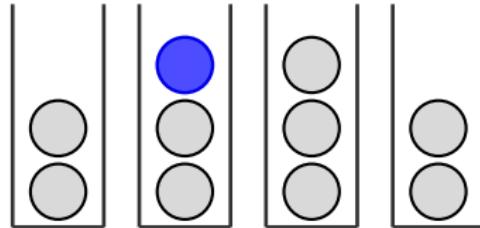
Computing Greedy($S \cup \{x^*\}$)

Two key observations:

ANALYZING THE RE COURSE



Computing **Greedy**(S)

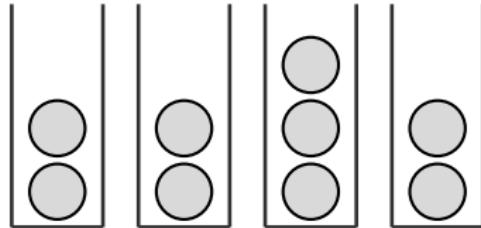


Computing **Greedy**($S \cup \{x^*\}$)

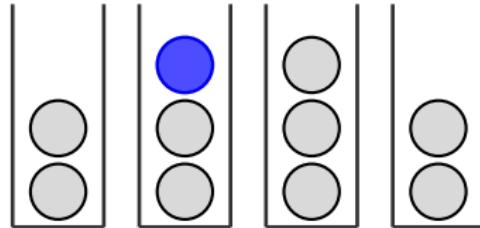
Two key observations:

1. There's always one special bin with an extra ball

ANALYZING THE RE COURSE



Computing **Greedy**(S)

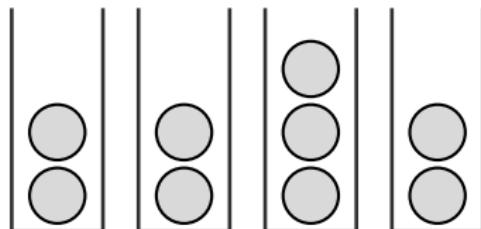


Computing **Greedy**($S \cup \{x^*\}$)

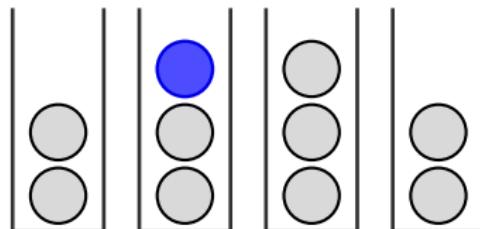
Two key observations:

1. There's always one special bin with an extra ball
2. If a ball incurs recourse, one of its choices is the special bin

ANALYZING THE RE COURSE



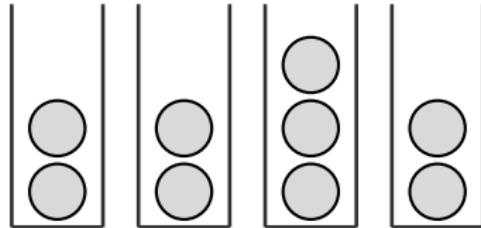
Computing **Greedy**(S)



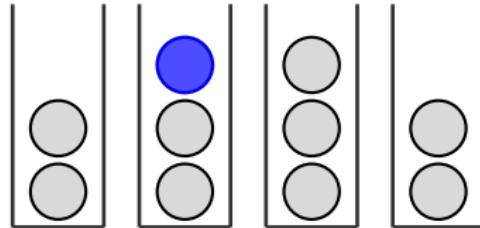
Computing **Greedy**($S \cup \{x^*\}$)

$$\Pr[\text{ball } x_i \text{ incurs recourse}] = O(1/n)$$

ANALYZING THE RECOURSE



Computing $\text{Greedy}(S)$



Computing $\text{Greedy}(S \cup \{x^*\})$

$$\Pr[\text{ball } x_i \text{ incurs recourse}] = O(1/n)$$

$$\implies \mathbb{E}[\text{total recourse}] = \sum_i \Pr[\text{ball } x_i \text{ incurs recourse}] = O(m/n)$$

THIS PAPER

Question: Does there exist a history-independent solution with small recourse and overload?

Theorem: There exists a history-independent solution with:

- ▶ Overload $O(1)$, with high probability.
- ▶ Expected recourse $O(\log \log(m/n))$.

Rest of Talk: A simple history-independent algorithm with overload $O(\log \log n)$ and expected recourse $O(m/n)$. ✓

History-Independent Load Balancing



Michael A. Bender
Stony Brook University



Bill Kuszmaul
CMU



Elaine Shi
CMU



Rose Silver
CMU