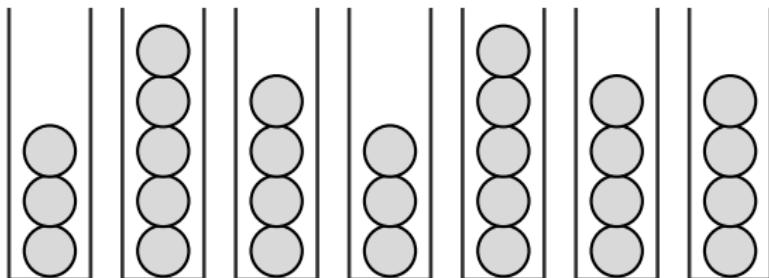


# History-Independent Load Balancing



Michael A. Bender

Stony Brook University

Bill Kuszmaul

CMU

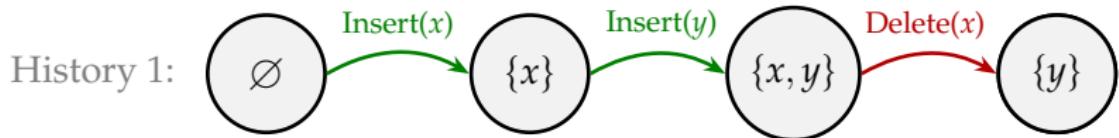
Elaine Shi

CMU

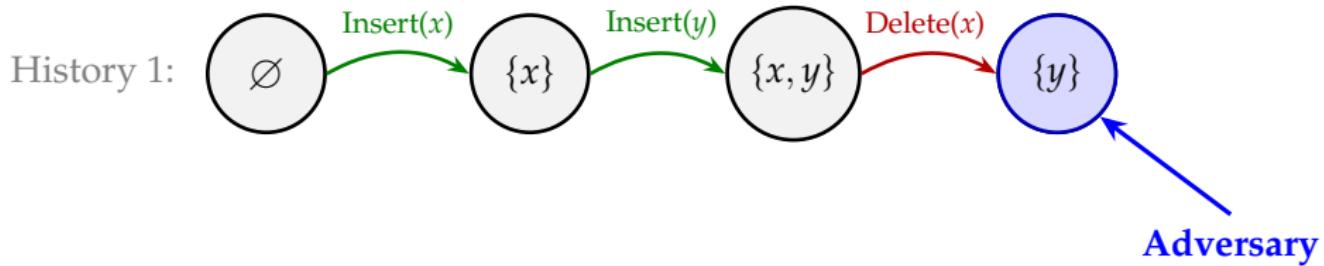
Rose Silver

CMU


# HISTORY-INDEPENDENT DATA STRUCTURES



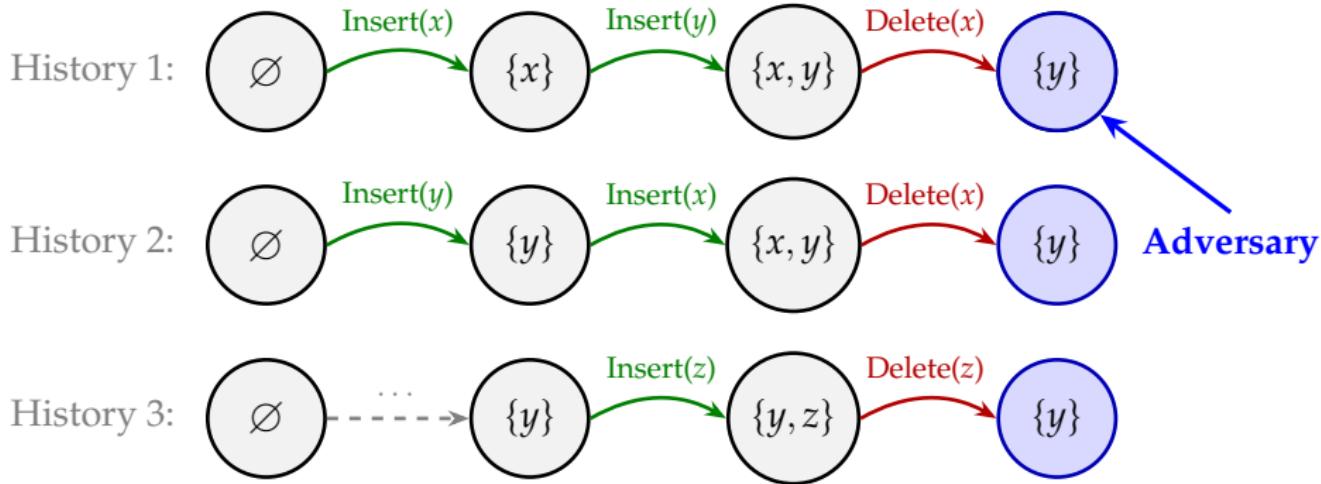
# HISTORY-INDEPENDENT DATA STRUCTURES



# HISTORY-INDEPENDENT DATA STRUCTURES



# HISTORY-INDEPENDENT DATA STRUCTURES



## History Independence (Micciancio '97, Naor & Teague '01)

- ▶ The state reveals only the current elements—not the history of operations.

# HISTORY INDEPENDENT DATA STRUCTURES

## A History of Applications

Hash tables, trees, memory allocation, PMAs, graph algorithms, cache-oblivious data structures, and more.

Micciancio '97, Naor & Teague '01, Buchbinder & Petrank '03, Molnar et al. '06, Blelloch & Golovin '07,  
Moran et al. '07, Naor et al. '08, Golovin '08–'10, Bajaj & Sion '13, Roche et al. '15, Bender et al. '16

# HISTORY INDEPENDENT DATA STRUCTURES

## A History of Applications

Hash tables, trees, memory allocation, PMAs, graph algorithms, cache-oblivious data structures, and more.

Micciancio '97, Naor & Teague '01, Buchbinder & Petrank '03, Molnar et al. '06, Blelloch & Golovin '07,  
Moran et al. '07, Naor et al. '08, Golovin '08–'10, Bajaj & Sion '13, Roche et al. '15, Bender et al. '16

**Yet fundamental questions remain open.**

# HISTORY INDEPENDENT DATA STRUCTURES

## A History of Applications


Hash tables, trees, memory allocation, PMAs, graph algorithms, cache-oblivious data structures, and more.

Micciancio '97, Naor & Teague '01, Buchbinder & Petrank '03, Molnar et al. '06, Blelloch & Golovin '07,  
Moran et al. '07, Naor et al. '08, Golovin '08–'10, Bajaj & Sion '13, Roche et al. '15, Bender et al. '16

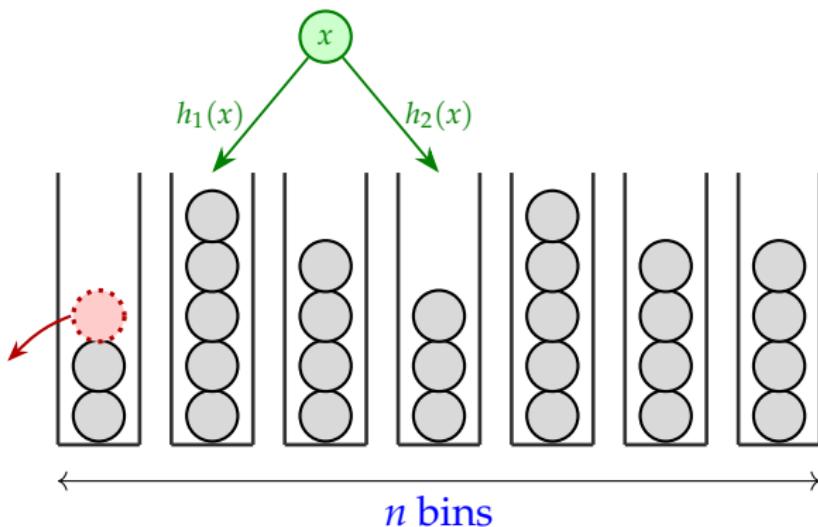
**Yet fundamental questions remain open.**

**This work:** History-Independent Load Balancing

# TWO-CHOICE LOAD BALANCING

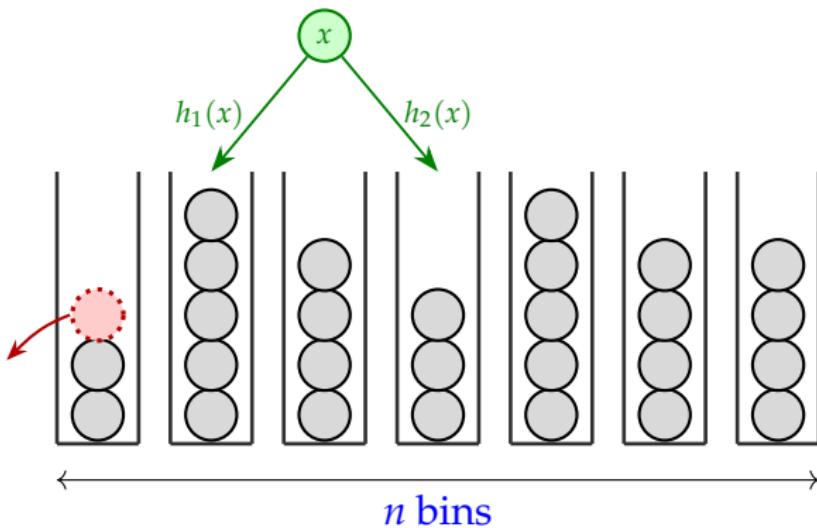


# TWO-CHOICE LOAD BALANCING




- ▶ Balls are **inserted/deleted**, with up to  $m$  present at a time.

# TWO-CHOICE LOAD BALANCING



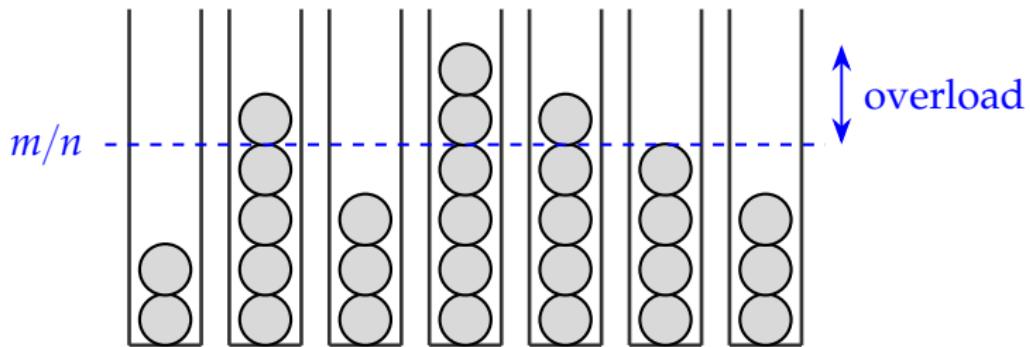
- ▶ Balls are **inserted/deleted**, with up to  $m$  present at a time.
- ▶ Each ball has two random bins where it can go.

# TWO-CHOICE LOAD BALANCING



- ▶ Balls are **inserted/deleted**, with up to  $m$  present at a time.
- ▶ Each ball has two random bins where it can go.
- ▶ We must maintain a valid assignment of balls to bins.

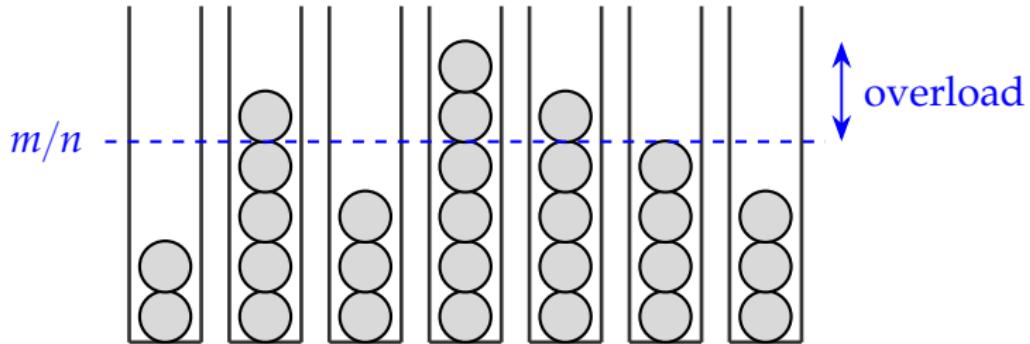
## TWO GOALS



### Minimize Overload:

- ▶ i.e., the amount by which the fullest bin exceeds  $m/n$ .

## TWO GOALS



### Minimize Overload:

- ▶ i.e., the amount by which the fullest bin exceeds  $m/n$ .

### Minimize Recourse:

- ▶ i.e., the number of balls moved around on any given insertion/deletion.

# THIS PAPER

**Question:** Does there exist a history-independent solution with small recourse and overload?

# THIS PAPER

**Question:** Does there exist a history-independent solution with small recourse and overload?

**Theorem:** There exists a history-independent solution with:

- ▶ Overload  $O(1)$ , with high probability.
- ▶ Expected recourse  $O(\log \log(m/n))$ .

# WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

# WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]  
[Frieze and Petti '18] ...

# WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]  
[Frieze and Petti '18] ...

But the fully dynamic case has remained largely open.

[Vöcking '99] [Dietzfelbinger and Weidling '07] [Bender, Conway, Farach-Colton, Kuszmaul, Tagliavini '21] [Bansal, Kuszmaul '22] ...

# WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]  
[Frieze and Petti '18] ...

But the fully dynamic case has remained largely open.

[Vöcking '99] [Dietzfelbinger and Weidling '07] [Bender, Conway, Farach-Colton, Kuszmaul, Tagliavini '21] [Bansal, Kuszmaul '22] ...

**Open Question:**

Is there a **fully dynamic** solution with **recourse  $o(m/n)$**  and **overload  $O(\log \log n)$** ?

# WHAT ABOUT NON-HISTORY-INDEPENDENT SOLUTIONS?

Lots of work on the insertion-only case.

[Azar, Broder, Karlin and Upfal '94] [Berenbrink, Czumaj, Steger, and Vöcking '00][Dietzfelbinger and Weidling '07]  
[Frieze and Petti '18] ...

But the fully dynamic case has remained largely open.

[Vöcking '99] [Dietzfelbinger and Weidling '07] [Bender, Conway, Farach-Colton, Kuszmaul, Tagliavini '21] [Bansal, Kuszmaul '22] ...

## Open Question:

Is there a **fully dynamic** solution with **recourse  $o(m/n)$**  and **overload  $O(\log \log n)$** ?

## Answer:

Yes! We get **recourse  $O(\log \log(m/n))$**  and **overload  $O(1)$** .

# THIS PAPER

**Question:** Does there exist a history-independent solution with small recourse and overload?

**Theorem:** There exists a history-independent solution with:

- ▶ Overload  $O(1)$ , with high probability.
- ▶ Expected recourse  $O(\log \log(m/n))$ .

# THIS PAPER

**Question:** Does there exist a history-independent solution with small recourse and overload?


**Theorem:** There exists a history-independent solution with:

- ▶ Overload  $O(1)$ , with high probability.
- ▶ Expected recourse  $O(\log \log(m/n))$ .

**Rest of Talk:** A simple history-independent algorithm with overload  $O(\log \log n)$  and expected recourse  $O(m/n)$ .


## WARMUP 1: THE SINGLE-CHOICE STRATEGY

To insert a ball  $x$ , just put it in bin  $h_1(x)$ :



## WARMUP 1: THE SINGLE-CHOICE STRATEGY

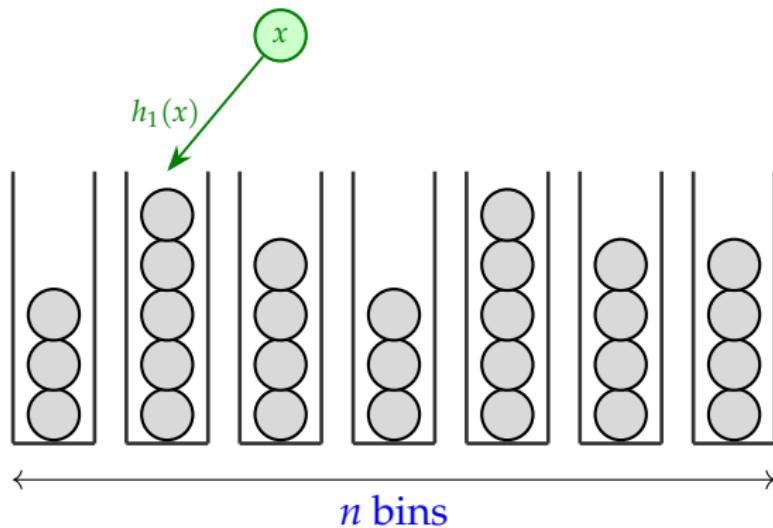
To insert a ball  $x$ , just put it in bin  $h_1(x)$ :



- ▶ This is history-independent ✓

## WARMUP 1: THE SINGLE-CHOICE STRATEGY

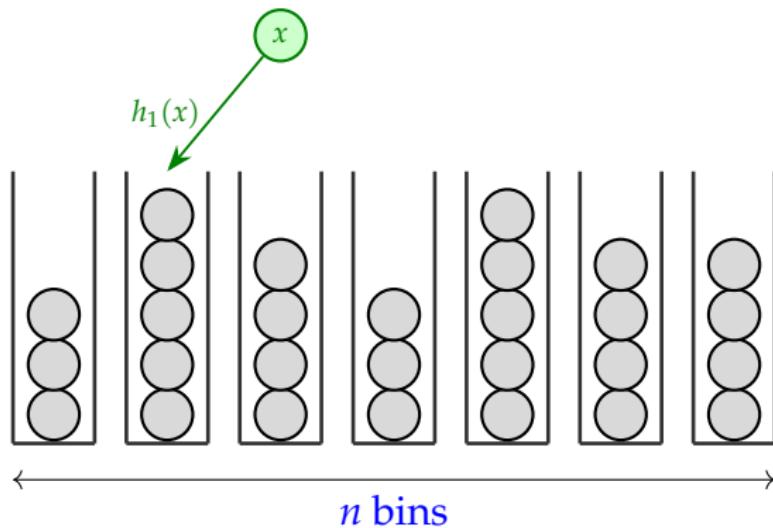
To insert a ball  $x$ , just put it in bin  $h_1(x)$ :



- ▶ This is history-independent ✓
- ▶ The recourse is 0 ✓

## WARMUP 1: THE SINGLE-CHOICE STRATEGY


To insert a ball  $x$ , just put it in bin  $h_1(x)$ :



- ▶ This is history-independent ✓
- ▶ The recourse is 0 ✓
- ▶ But... the overload is huge, roughly  $\sqrt{m/n}$  ✗

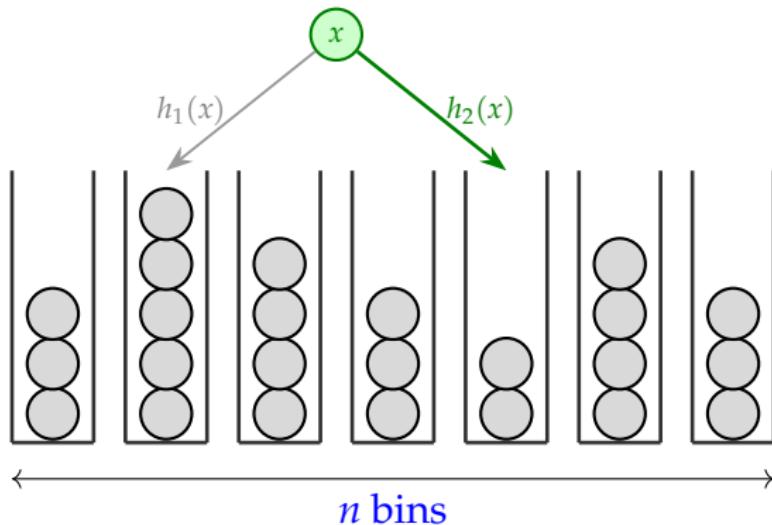
## WARMUP 2: GREEDY INSERTIONS

To insert a ball  $x$ , put it in the **emptier** of its choices:



## WARMUP 2: GREEDY INSERTIONS


To insert a ball  $x$ , put it in the **emptier** of its choices:



- ▶ This is **not** history-independent  $\times$

## WARMUP 2: GREEDY INSERTIONS


To insert a ball  $x$ , put it in the **emptier** of its choices:



- ▶ This is **not** history-independent ✗
- ▶ The recourse is 0 ✓

## WARMUP 2: GREEDY INSERTIONS

To insert a ball  $x$ , put it in the **emptier** of its choices:

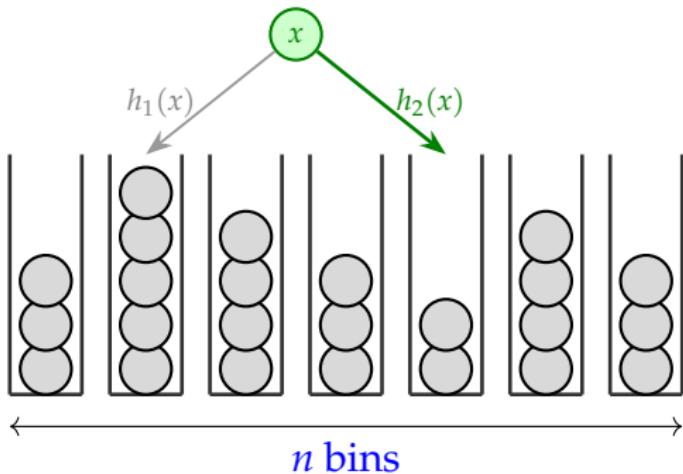


- ▶ This is **not** history-independent ✗
- ▶ The recourse is 0 ✓
- ▶ In the insertion-only case, the overload is  $O(\log \log n)$  ✓

[Azar, Broder, Karlin and Upfal '94]

# A SIMPLE HISTORY-INDEPENDENT ALGORITHM


# A SIMPLE HISTORY-INDEPENDENT ALGORITHM



Given a set  $S$  of balls, define  $\text{Greedy}(S)$  as:

- ▶ Start with empty bins.
- ▶ Sort the balls in  $S$  to get a sequence  $x_1, x_2, \dots$ .
- ▶ Insert  $x_1, x_2, \dots$  using the greedy algorithm.

# A SIMPLE HISTORY-INDEPENDENT ALGORITHM




Given a set  $S$  of balls, define  $\text{Greedy}(S)$  as:


- ▶ Start with empty bins.
- ▶ Sort the balls in  $S$  to get a sequence  $x_1, x_2, \dots$ .
- ▶ Insert  $x_1, x_2, \dots$  using the greedy algorithm.

**A History-Independent Algorithm:** If  $S$  is the current set, use  $\text{Greedy}(S)$ .

# ANALYZING HISTORY-INDEPENDENT GREEDY




# ANALYZING HISTORY-INDEPENDENT GREEDY




- ▶ The algorithm is history independent ✓

# ANALYZING HISTORY-INDEPENDENT GREEDY



# ANALYZING HISTORY-INDEPENDENT GREEDY



- ▶ The algorithm is history independent ✓
- ▶ The overload is  $O(\log \log n)$  ✓
- ▶ What is the recourse?

## ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

How does  $\text{Greedy}(S)$  change if we add a ball  $x^*$ ?

## ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )



Computing **Greedy**( $S \cup \{x^*\}$ )

# ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

# ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

# ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

## ANALYZING THE RE COURSE




Computing  $\text{Greedy}(S)$




Computing  $\text{Greedy}(S \cup \{x^*\})$

# ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

## ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

# ANALYZING THE RE COURSE


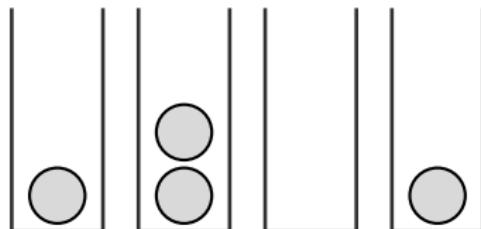


Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )


# ANALYZING THE RE COURSE



# ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

## ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

Subsequent balls will experience either:

## ANALYZING THE RE COURSE



**Computing Greedy( $S$ )**




**Computing Greedy( $S \cup \{x^*\}$ )**


Subsequent balls will experience either:

1. No recourse

## ANALYZING THE RE COURSE



Computing Greedy( $S$ )




Computing Greedy( $S \cup \{x^*\}$ )


Future insertions will experience either:

1. No recourse

## ANALYZING THE RE COURSE



Computing Greedy( $S$ )




Computing Greedy( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse

## ANALYZING THE RE COURSE



Computing Greedy( $S$ )




Computing Greedy( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse

# ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse

## ANALYZING THE RE COURSE



Computing Greedy( $S$ )





Computing Greedy( $S \cup \{x^*\}$ )

Subsequent balls will experience either:

1. No recourse


## ANALYZING THE RE COURSE




Subsequent balls will experience either:

1. No recourse

## ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )





Computing **Greedy**( $S \cup \{x^*\}$ )

Subsequent balls will experience either:

1. No recourse
2. Recourse


## ANALYZING THE RE COURSE




Subsequent balls will experience either:

1. No recourse
2. Recourse

# ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse
2. Recourse

## ANALYZING THE RE COURSE



Computing Greedy( $S$ )




Computing Greedy( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse
2. Recourse

## ANALYZING THE RE COURSE



Computing Greedy( $S$ )




Computing Greedy( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse
2. Recourse

# ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )


Subsequent balls will experience either:

1. No recourse
2. Recourse

## ANALYZING THE RE COURSE



**Computing Greedy( $S$ )**




**Computing Greedy( $S \cup \{x^*\}$ )**


Subsequent balls will experience either:

1. No recourse
2. Recourse

# ANALYZING THE RE COURSE




**Computing Greedy( $S$ )**




**Computing Greedy( $S \cup \{x^*\}$ )**

Two key observations:

## ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )


Two key observations:

1. There's always one special bin with an extra ball

## ANALYZING THE RE COURSE



Computing **Greedy**( $S$ )




Computing **Greedy**( $S \cup \{x^*\}$ )

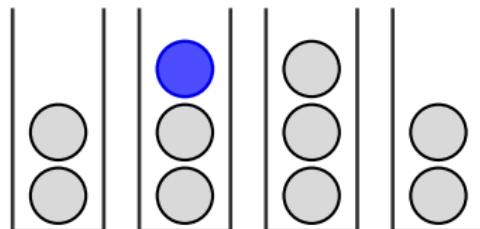
Two key observations:

1. There's always one special bin with an extra ball
2. If a ball incurs recourse, one of its choices is the special bin

# ANALYZING THE RE COURSE




Computing **Greedy**( $S$ )



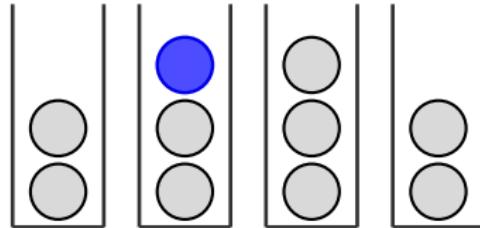
Computing **Greedy**( $S \cup \{x^*\}$ )

$$\Pr[\text{ball } x_i \text{ incurs recourse}] = O(1/n)$$

# ANALYZING THE RECOURSE



**Computing  $\text{Greedy}(S)$**



**Computing  $\text{Greedy}(S \cup \{x^*\})$**

$$\Pr[\text{ball } x_i \text{ incurs recourse}] = O(1/n)$$

$$\implies \mathbb{E}[\text{total recourse}] = \sum_i \Pr[\text{ball } x_i \text{ incurs recourse}] = O(m/n)$$

# THIS PAPER


**Question:** Does there exist a history-independent solution with small recourse and overload?

**Theorem:** There exists a history-independent solution with:

- ▶ Overload  $O(1)$ , with high probability.
- ▶ Expected recourse  $O(\log \log(m/n))$ .

**Rest of Talk:** A simple history-independent algorithm with overload  $O(\log \log n)$  and expected recourse  $O(m/n)$ . ✓

# History-Independent Load Balancing



Michael A. Bender

Stony Brook University

Bill Kuszmaul

CMU

Elaine Shi

CMU

Rose Silver

CMU